

COMMON METRICS

INTRODUCTION

As a data scientist, when you're not investigating spikes or dips in your data, you might be building dashboards of *KPIs*, or *key performance indicators* for a company.

KPIs are often displayed on TVs on the walls of the office, and serve as high level health metrics for the business. While every company's metrics are defined slightly differently, the basics are usually very similar.

In this lesson we'll take a look at basic KPIs like Daily Revenue, Daily Active Users, ARPU, and Retention for a video game, **Mineblocks**.

1. This company has two tables, `gameplays` and `purchases`. The `purchases` table lists all purchases made by players while they're playing Mineblocks. Complete the query to select from `purchases`.

```
select *  
from /**/  
order by id  
limit 10;
```

```
select *  
from purchases  
order by id  
limit 10;
```

2. The `gameplays` table lists the date and platform for each session a user plays. Select from `gameplays`.

```
select *  
from /**/  
order by id  
limit 10;
```

```
select *  
from gameplays  
order by id  
limit 10;
```

DAILY REVENUE

At the heart of every company is revenue, and Mineblocks is no exception. For our first KPI we'll calculate daily revenue.

1. Daily Revenue is simply the sum of money made per day. To get close to Daily Revenue, we calculate the daily sum of the prices in the `purchases` table. Complete the query by using the `sum` function and passing the `price` column from the `purchases` table.

```
select
  date(created_at),
  round(sum(**/), 2)
from **/
group by 1
order by 1;
```

```
select date(created_at), round(sum(price), 2)
from purchases
group by 1
order by 1
limit 10;
```

DAILY REVENUE 2

Great! That query doesn't take refunds into account. We'll update the query to exclude refunds. Fields like `refunded_at` will only have data if the transaction was refunded, and otherwise left `null`.

1. Update our daily revenue query to exclude refunds. Complete the query by filtering for `refunded_at` is `null`.

```
select
  date(created_at),
  round(sum(price), 2) as daily_rev
from purchases
where **/
group by 1
order by 1;
```

```
select date(created_at), round(sum(price), 2) as daily_rev
from purchases
where refunded_at is null
group by 1
order by 1
limit 10;
```

DAILY ACTIVE USERS

Mineblocks is a game, and one of the core metrics to any game is the number of people who play each day. That KPI is called *Daily Active Users*, or *DAU*.

DAU is defined as the number of unique players seen in-game each day. It's important not to double count users who played multiple times, so we'll use `distinct` in our `count` function.

Likewise, Weekly Active Users (WAU) and Monthly Active Users (MAU) are in the same family.

1. For Mineblocks, we'll use the `gameplays` table to calculate DAU. Each time a user plays the game, their session is recorded in `gameplays`. Thus, a distinct count of users per day from `gameplays` will give us DAU.

Calculate Daily Active Users for Mineblocks. Complete the query's `count` function by passing in the `distinct` keyword and the `user_id` column name.

```
select
  date(created_at),
  count(/**/) as dau
from gameplays
group by 1
order by 1;
```

```
select date(created_at), count(distinct user_id) as dau
from gameplays
group by 1
order by 1
limit 10;
```

Here's a [hint](#) on how to use the count function to count distinct columns in a table. Also notice that we include `platform` in both the `select` and `group by` clauses.

DAILY ACTIVE USERS 2

Great! Since Mineblocks is on multiple platforms, we can calculate DAU per-platform.

1. Previously we calculated DAU only per day, so the output we wanted was `[date, dau_count]`. Now we want DAU per platform, making the desired output `[date, platform, dau_count]`.

Calculate DAU for Mineblocks per-platform. Complete the query below. You will need to select the `platform` column and add a `count` function by passing in the `distinct` keyword and the `user_id` column name.

```
select
  date(created_at),
  /**/,
  count(/**/) as dau
from gameplays
group by 1, 2
order by 1, 2;
```

```
select date(created_at), platform, count(distinct user_id) as dau
from gameplays
group by 1, 2
order by 1, 2
limit 10;
```

DAILY AVERAGE REVENUE PER PURCHASING USER

We've looked at DAU and Daily Revenue in Mineblocks. Now we must understand the purchasing habits of our users.

Mineblocks, like every freemium game, has two types of users:

- purchasers: users who have bought things in the game
- players: users who play the game but have not yet purchased

The next KPI we'll look at **Daily ARPPU** - Average Revenue Per Purchasing User. This metric shows if the average amount of money spent by purchasers is going up over time.

Daily ARPPU is defined as the sum of revenue divided by the number of purchasers per day.

1. To get Daily ARPPU, modify the daily revenue query from earlier to divide by the number of purchasers.

Complete the query by adding a numerator and a denominator. The numerator will display daily revenue, or `sum` the `price` columns. The denominator will display the number of purchasers by passing the `distinct` keyword and the `user_id` column name into the `count` function.

```
select
  date(created_at),
  round(**/ / count(**/), 2) as arppu
from purchases
where refunded_at is null
group by 1
order by 1;
```

```
select date(created_at), round(sum(price) / count(distinct(user_id)), 2) as arppu
from purchases
where refunded_at is null
group by 1
order by 1
limit 10;
```

DAILY AVERAGE REVENUE PER USER

The more popular (and difficult) cousin to Daily ARPPU is *Daily ARPU*, Average Revenue Per User. ARPU measures the average amount of money we're getting across all players, whether or not they've purchased.

ARPPU increases if purchasers are spending more money. ARPU increases if more players are choosing to purchase, even if the purchase size stays consistent.

No one metric can tell the whole story. That's why it's so helpful to have many KPIs on the same dashboard.

ARPU 2

Daily ARPU is defined as revenue divided by the number of players, per-day. To get that, we'll need to calculate the daily revenue and daily active users separately, and then join them on their dates.

One way to easily create and organize temporary results in a query is with *CTEs, Common Table Expressions*, also known as **with** clauses. The **with** clauses make it easy to define and use results in a more organized way than subqueries.

These clauses usually look like this:

```
with {subquery_name} as ( {subquery_body} )
select ... from {subquery_name} where ...
```

1. Use a **with** clause to define **daily_revenue** and then select from it.

```
/**/ daily_revenue as (
  select
    date(created_at) as dt,
    round(sum(price), 2) as rev
  from purchases
  where refunded_at is null
  group by 1
)
select * from daily_revenue order by dt;
```

```
with daily_revenue as (
  select date(created_at) as dt, round(sum(price), 2) as rev
  from purchases
  where refunded_at is null
  group by 1
)
select * from daily_revenue order by dt;
```

ARPU 2

Great! Now you're familiar with using the **with** clause to create temporary result sets.

You just built the first part of ARPU, **daily_revenue**. From here we can build the second half of ARPU in our **with** clause, **daily_players**, and use both together to create ARPU.

1. Building on this CTE, we can add in DAU from earlier. Complete the query by calling the DAU query we created earlier, now aliased as **daily_players**:

```
/**/ daily_revenue as (
  select
    date(created_at) as dt,
    round(sum(price), 2) as rev
  from purchases
  where refunded_at is null
  group by 1
```

```

),
daily_players as (
  select
    /**/ as dt,
    /**/ as players
  from gameplays
  group by 1
)
select * from daily_players order by dt;

```

```

with daily_revenue as (
  select date(created_at) as dt, round(sum(price), 2) as rev
  from purchases
  where refunded_at is null
  group by 1
),
daily_players as (
  select
    date(created_at) as dt, count(distinct user_id) as players
  from gameplays
  group by 1
)
select * from daily_players order by dt
limit 10;

```

Here's a [hint](#) on how we created the previous DAU query.

- Now that we have the revenue and DAU, join them on their dates and calculate daily ARPU. Complete the query by adding the keyword `using` in the `join` clause.

```

/**/ daily_revenue as (
  select
    date(created_at) as dt,
    round(sum(price), 2) as rev
  from purchases
  where refunded_at is null
  group by 1
),
daily_players as (
  select
    /**/ as dt,
    /**/ as players
  from gameplays
  group by 1
)
select
  daily_revenue.dt,
  daily_revenue.rev / daily_players.players
from daily_revenue
join daily_players /**/ (dt);

```

```

with daily_revenue as (
select date(created_at) as dt, round(sum(price), 2) as rev
from purchases
where refunded_at is null
group by 1
),
daily_players as (
select
date(created_at) as dt, count(distinct user_id) as players
from gameplays
group by 1
)
select daily_revenue.dt, daily_revenue.rev / daily_players.players
from daily_revenue
join daily_players using (dt)
limit 10;

```

In the final `select` statement, `daily_revenue.dt` represents the date, while `daily_revenue.rev / daily_players.players` is the daily revenue divided by the number of players that day. In full, it represents how much the company is making per player, per day.

ARPU 3

Nice work, you just defined ARPU for Mineblocks! In our ARPU query, we used `using` instead of `on` in the `join` clause. This is a special case join.

```

from daily_revenue
  join daily_players using (dt);

```

When the columns to join have the same name in both tables you can use `using` instead of `on`. Our use of the `using` keyword is in this case equivalent to this clause:

```

from daily_revenue
  join daily_players on
    daily_revenue.dt = daily_players.dt;

```

1 DAY RETENTION

Now let's find out what percent of Mineblock players are returning to play the next day. This KPI is called *1 Day Retention*.

Retention can be defined many different ways, but we'll stick to the most basic definition. For all players on Day N, we'll consider them retained if they came back to play again on Day N+1.

This will let us track whether or not Mineblocks is getting "stickier" over time. The stickier our game, the more days players will spend in-game.

And more time in-game means more opportunities to monetize and grow our business.

1 DAY RETENTION 2

Before we can calculate retention we need to get our data formatted in a way where we can determine if a user returned.

Currently the `gameplays` table is a list of when the user played, and it's not easy to see if any user came back.

By using a `self-join`, we can make multiple gameplays available on the same row of results. This will enable us to calculate retention.

The power of `self-join` comes from joining every row to every other row. This makes it possible to compare values from two different rows in the new result set. In our case, we'll compare rows that are one date apart from each user.

1. To calculate retention, start from a query that selects the `date(created_at)` as `dt` and `user_id` columns from the `gameplays` table.

```
select
  date(/**/) as dt,
  /**/
from gameplays as g1
order by dt
limit 100;
```

```
select date(created_at) as dt, user_id
from gameplays as g1
order by dt
limit 100;
```

2. Now we'll join `gameplays` on itself so that we can have access to all gameplays for each player, for each of their gameplays.

This is known as a `self-join` and will let us connect the players on Day N to the players on Day N+1. In order to join a table to itself, it must be aliased so we can access each copy separately.

We aliased `gameplays` in the query above because in the next step, we need to join `gameplays` to itself so we can get a result selecting `[date, user_id, user_id_if_retained]`.

Complete the query by using a `join` statement to join `gameplays` to itself on `user_id` using the aliases `g1` and `g2`.

```
select
  date(g1.created_at) as dt,
  g1.user_id
from gameplays as g1
```



```

/**/ gameplays as g2 on
  g1.user_id = g2.user_id
order by 1
limit 100;

```

```

select date(g1.created_at) as dt, g1.user_id
from gameplays as g1
join gameplays as g2
on g1.user_id = g2.user_id
order by 1
limit 100;

```

We don't use the `using` clause here because the `join` is about to get more complicated.

1 DAY RETENTION 3

Now that we have our `gameplays` table joined to itself, we can start to calculate retention.

1 Day Retention is defined as the number of players who returned the next day divided by the number of original players, per day. Suppose 10 players played Mineblocks on Dec 10th. If 4 of them play on Dec 11th, the 1 day retention for Dec 10th is 40%.

1. The previous query joined all rows in `gameplays` against all other rows for each user, making a massive result set that we don't need. We'll need to modify this query.

```

select
  date(g1.created_at) as dt,
  g1.user_id,
  g2.user_id
from gameplays as g1
join gameplays as g2 on
  g1.user_id = g2.user_id
  and /**/
order by 1
limit 100;

```

Complete the query above such that the `join` clause includes a date join:

```

date(g1.created_at) = date(datetime(g2.created_at, '-1 day'))

```

```

select date(g1.created_at) as dt, g1.user_id, g2.user_id
from gameplays as g1
join gameplays as g2
on g1.user_id = g2.user_id
and date(g1.created_at) = date(datetime(g2.created_at, '-1 day'))
order by 1
limit 100;

```

This means "only join rows where the date in `g1` is one less than the date in `g2`", which makes it possible to see if users have returned!

2. The query above won't return meaningful results because we're using an `inner join`. This type of join requires that the condition be met for all rows, effectively limiting our selection to only the users that have returned.

Instead, we want to use a `left join`, this way all rows in `g1` are preserved, leaving nulls in the rows from `g2` where users did not return to play the next day.

Change the `join` clause to use `left join` and count the distinct number of users from `g1` and `g2` per date.

```
select
  date(g1.created_at) as dt,
  count(distinct g1.user_id) as total_users,
  count(distinct g2.user_id) as retained_users
from gameplays as g1
  /**/ gameplays as g2 on
  g1.user_id = g2.user_id
  and date(g1.created_at) = date(datetime(g2.created_at, '-1 day'))
group by 1
order by 1
limit 100;
```

```
select
  date(g1.created_at) as dt,
  count(distinct g1.user_id) as total_users,
  count(distinct g2.user_id) as retained_users
from gameplays as g1
left join gameplays as g2 on
  g1.user_id = g2.user_id
  and date(g1.created_at) = date(datetime(g2.created_at, '-1 day'))
group by 1
order by 1
```

3. Now that we have retained users as `count(distinct g2.user_id)` and total users as `count(distinct g1.user_id)`, divide retained users by total users to calculate 1 day retention!

```
select
  date(g1.created_at) as dt,
  round(100 * count(**/) /
    count(**/)) as retention
from gameplays as g1
  left join gameplays as g2 on
  g1.user_id = g2.user_id
  and date(g1.created_at) = date(datetime(g2.created_at, '-1 day'))
group by 1
order by 1
limit 100;
```

```
select
  date(g1.created_at) as dt,
  round(100 * count(distinct g2.user_id) /
        count(distinct g1.user_id)) as retention
from gameplays as g1
  left join gameplays as g2 on
    g1.user_id = g2.user_id
  and date(g1.created_at) = date(datetime(g2.created_at, '-1 day'))
group by 1
order by 1
limit 100;
```

COMMON METRICS CONCLUSION

While every business has different metrics to track their success, most are based on revenue and usage.

The metrics in this lesson are merely a starting point, and from here you'll be able to create and customize metrics to track whatever is most important to your company.

And remember, data science is exploratory! The current set of metrics can always be improved and there's usually more to any spike or dip than what immediately meets the eye.

Let's generalize what we've learned so far:

- *Key Performance Indicators* are high level health metrics for a business.
- *Daily Revenue* is the sum of money made per day.
- *Daily Active Users* are the number of unique users seen each day
- *Daily Average Revenue Per Purchasing User (ARPPU)* is the average amount of money spent by purchasers each day.
- *Daily Average Revenue Per User (ARPU)* is the average amount of money across all users.
- *1 Day Retention* is defined as the number of players from Day N who came back to play again on Day N+1.